

# Multithreading

# Parallel Processing

- Some programs can run faster if tasks can be run in parallel
  - Calculations on large sets of numbers if operations can be done in parallel
    - $(1+3)*(2+7)$ : additions can be done in parallel, then add results
  - Animation work calculates positions in parallel

# Processes and Threads

- **Process**

- A self contained execution environment
  - Every process has its own memory, cpu time, etc
- Processes cannot share memory between themselves directly

- **Thread**

- A program unit that runs independently of other parts of the program
- Can have multiple threads in a single process
- Can share memory with other threads in the same process

- Usually, the JVM executes each thread for a short period of time, then switches
  - This appears to run in parallel, but isn't
  - This works well; IO and network ops don't require the CPU
- Multiple processor/core computers allow multiple threads to run simultaneously

- To run a new thread in Java
  - In the class that represents your thread
    - define a class that implements the `Runnable` interface
    - override the `run` method of the interface
  - In your driver
    - instantiate an object of your class
    - construct a `Thread` object from the object you instantiated
    - call the `start` method of the `Thread` object

- See the included threading project for an example of basic threading

# Terminating Threads

- A thread terminates when the `run` finishes
- To stop a thread manually, you should `interrupt` it
  - `stop` is deprecated, do not use it
  - When the `interrupt` method is called, a `boolean` field in the thread data structure is set
  - The `run` method can check for interruptions, do required cleanup, then exit

# Dangers of Threading

- Race Conditions
  - Two threads compete for the same resources
  - Unpredictable results (who gets the resources first?)
  - Fixed using locks
- Locks not released
  - Fix using `finally`
- Deadlock
  - Fix using `conditions`



# Race Condition Example

- Two threads are updating a bank account (one is withdrawing \$100, one is depositing \$100)
  - initial balance is 0
  - thread 1 reads initial balance of 0
  - thread 2 reads initial balance of 0
  - thread 1 deposits 100 and calculates (but doesn't store) the new balance of 100
  - thread 2 withdraws 100 and calculates (but doesn't store) the new balance of -100
  - thread 2 stores the new balance of -100
  - thread 1 stores the new balance of 100
  - final balance: 100! (could have been -100 if the thread order had been different)

- Lock objects can be used to solve race conditions
- **ReentrantLock** class
  - Use this with a class whose methods access shared resources  
`balanceChangeLock = new ReentrantLock( );`
  - lock and unlock the `ReentrantLock` as required  
`balanceChangeLock.lock( );`  
`// code that uses shared resources`  
`balanceChangeLock.unlock( );`
  - Only one thread may have the object locked at a time; others must wait until it is unlocked again

# Handling Exceptions

- If an exception is thrown between calls to `lock` and `unlock`, the call to `unlock` never happens
- To solve this, use a `try` block around the code using the shared resource, and place the call to `unlock` in a `finally` block

- We want to add a test to disallow negative balances
  - Add a test inside the try block: if the balance is less than the amount being withdrawn, we wait until the balance grows
- How do we wait?
  - If the thread sleeps, it will block other threads trying to acquire the same lock it's holding
  - Other threads will call deposit, but be blocked by the withdraw method
  - Withdraw can't complete until a deposit is made!

# Condition Objects

- Conditions allow an object to temporarily release a lock and regain it later
- A condition belongs to a lock object
- The `newCondition` method allows a lock to obtain a condition
- You can call the `await` method of the `Condition` object to make the thread holding the lock wait and allow another thread to acquire the lock objects
- The thread with the condition is blocked until another thread executes the `signalAll` method on the same condition object

# Synchronized Methods

- Earlier Java versions used synchronized methods
- Every Java object has one built-in lock and one built-in condition
- if a synchronized method is called, the lock is automatically created
- You can make any method synchronized by adding the `synchronized` keyword to the method header  
`public synchronized void deposit (double amount)`
- Synchronized methods automatically get the lock/try/finally/unlock code implemented for it
- Conditions also work by using the `wait()` method  
`while(balance < amount) wait();`
  - Another method must call `notifyAll()` to end the wait
- Synchronized methods are easier, but more limited
  - Use these if they will work for your purpose